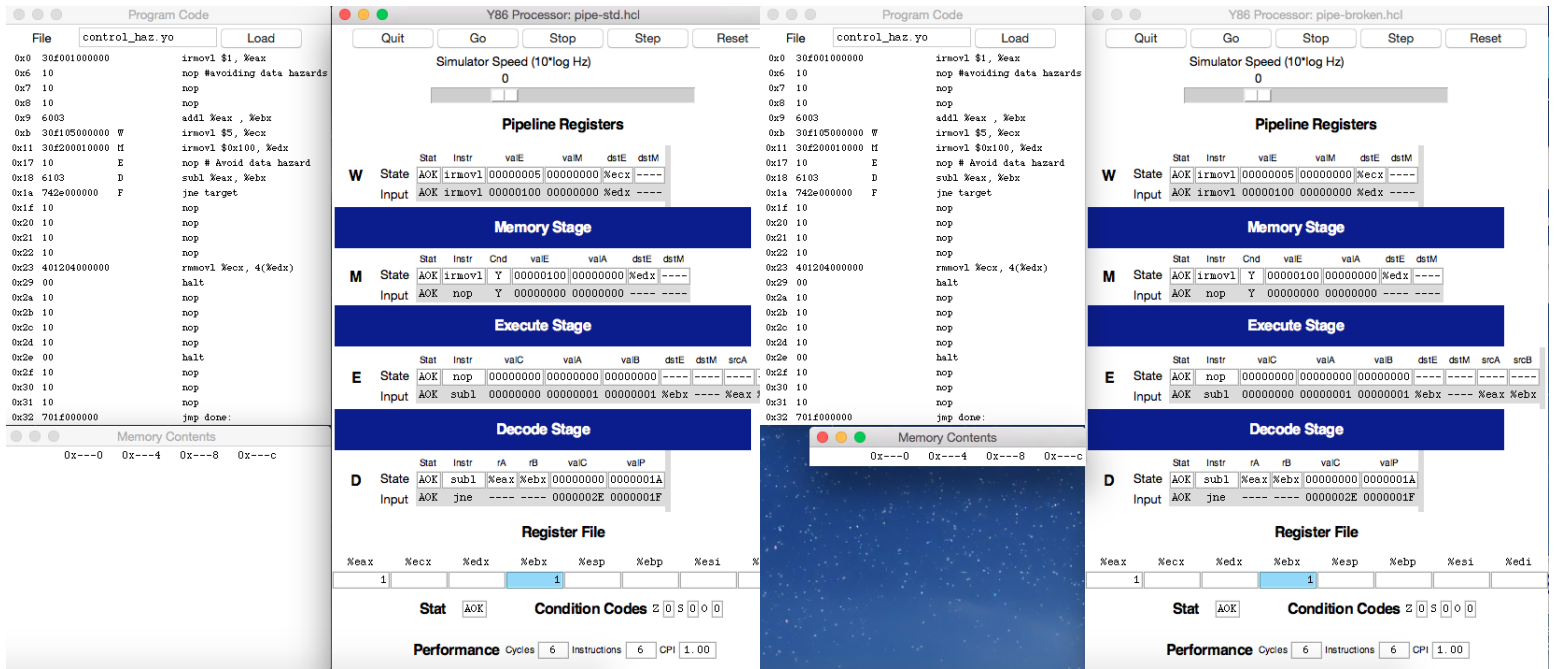


# Control Hazard Demonstration

Carter J. Bastian \*

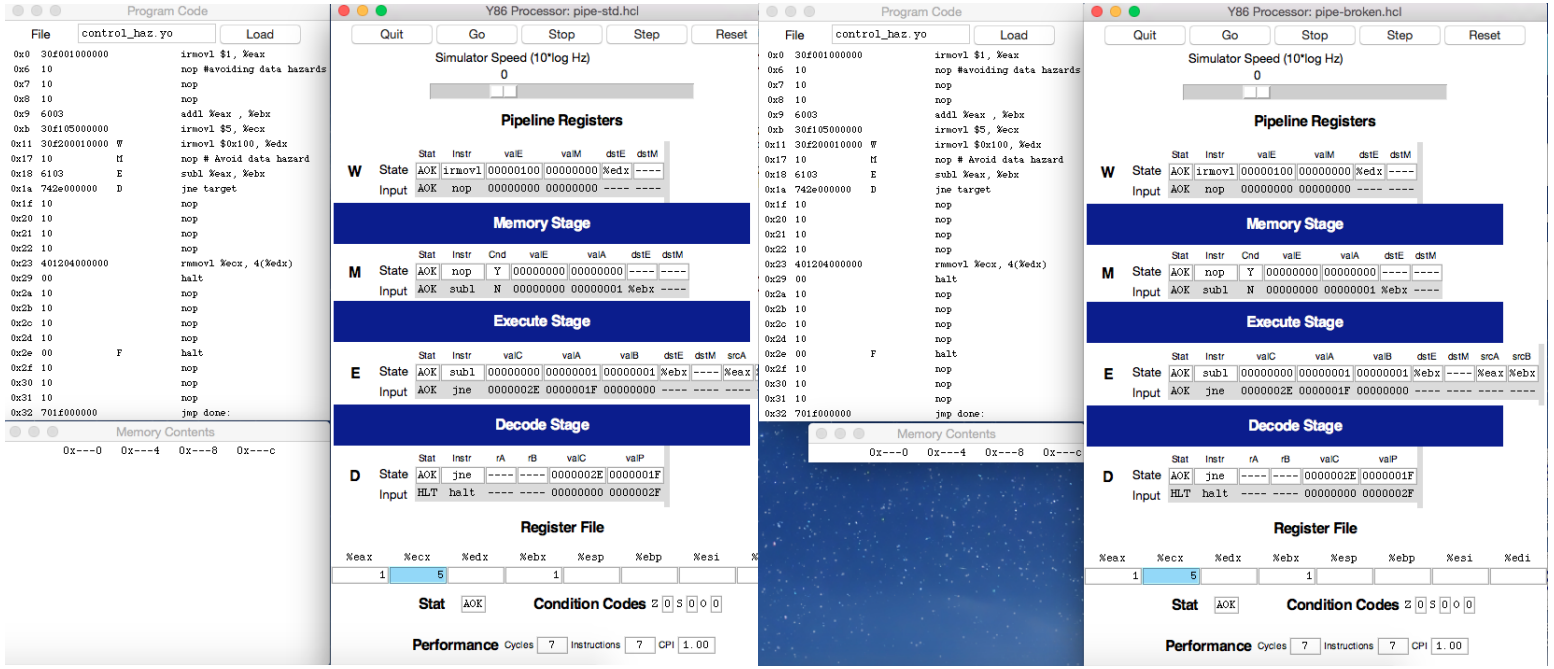
May 2015



(a) Correct Implementation (std)

(b) Incorrect Implementation (broken)

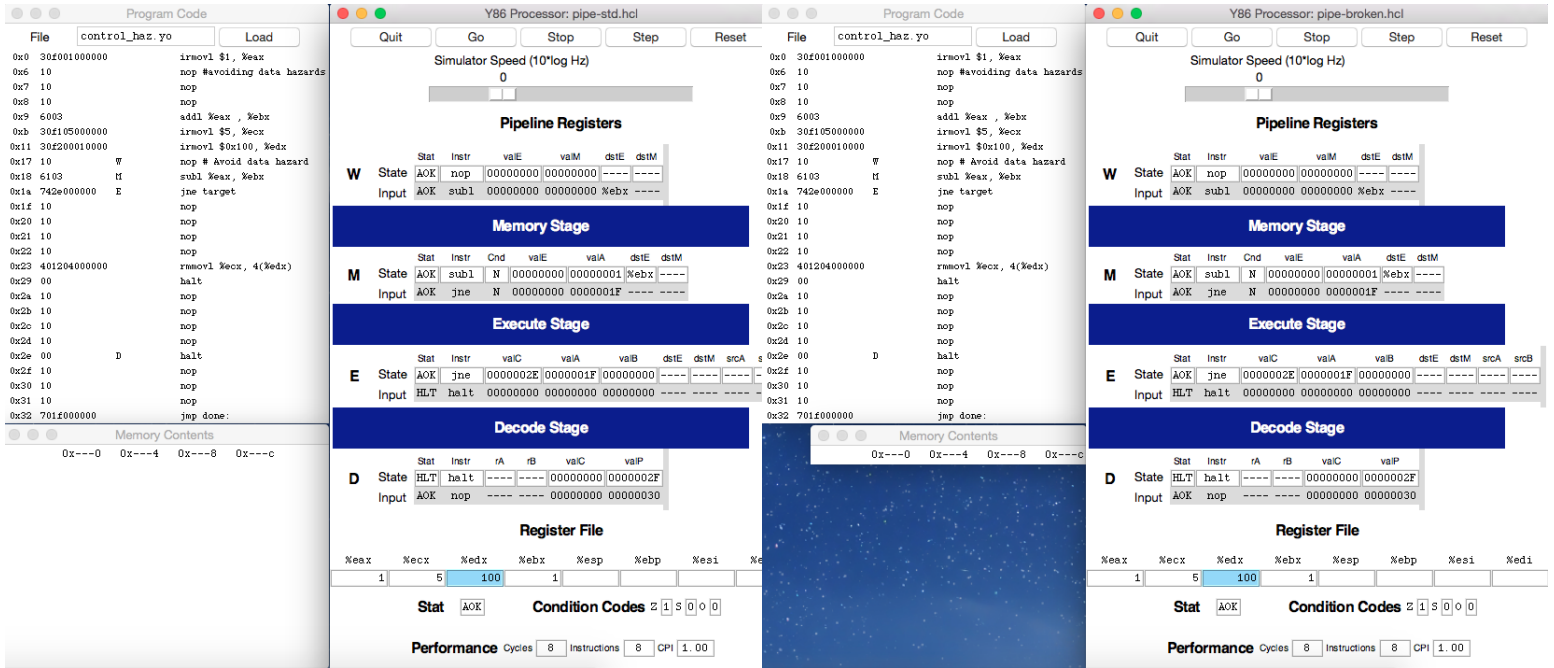
Figure 1: In this step, the `jne` instruction has just been fetched by the two different versions.



(a) Correct Implementation (std)

(b) Incorrect Implementation (broken)

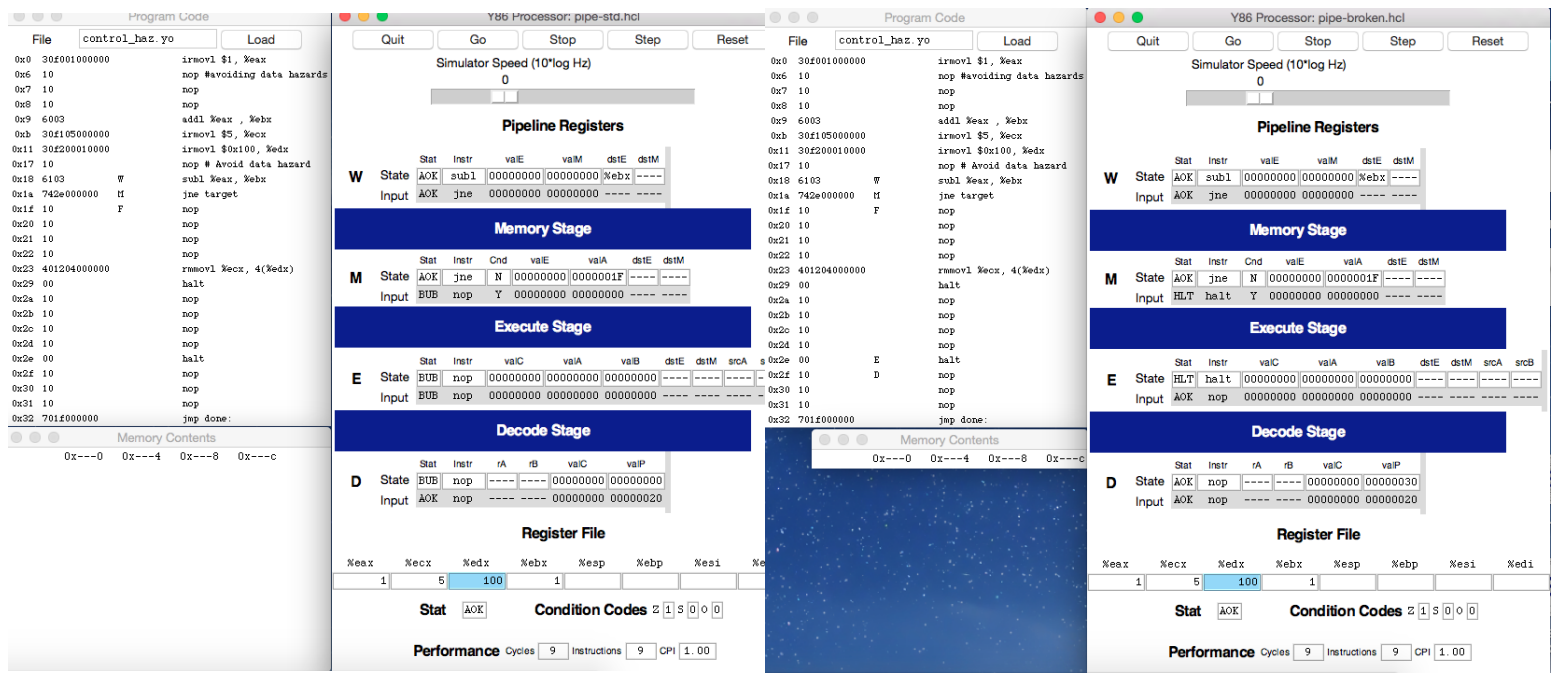
Figure 2: In this step, both implementations predict the conditional jump to “target” and fetch the `halt` instruction which should not be executed.



(a) Correct Implementation (std)

(b) Incorrect Implementation (broken)

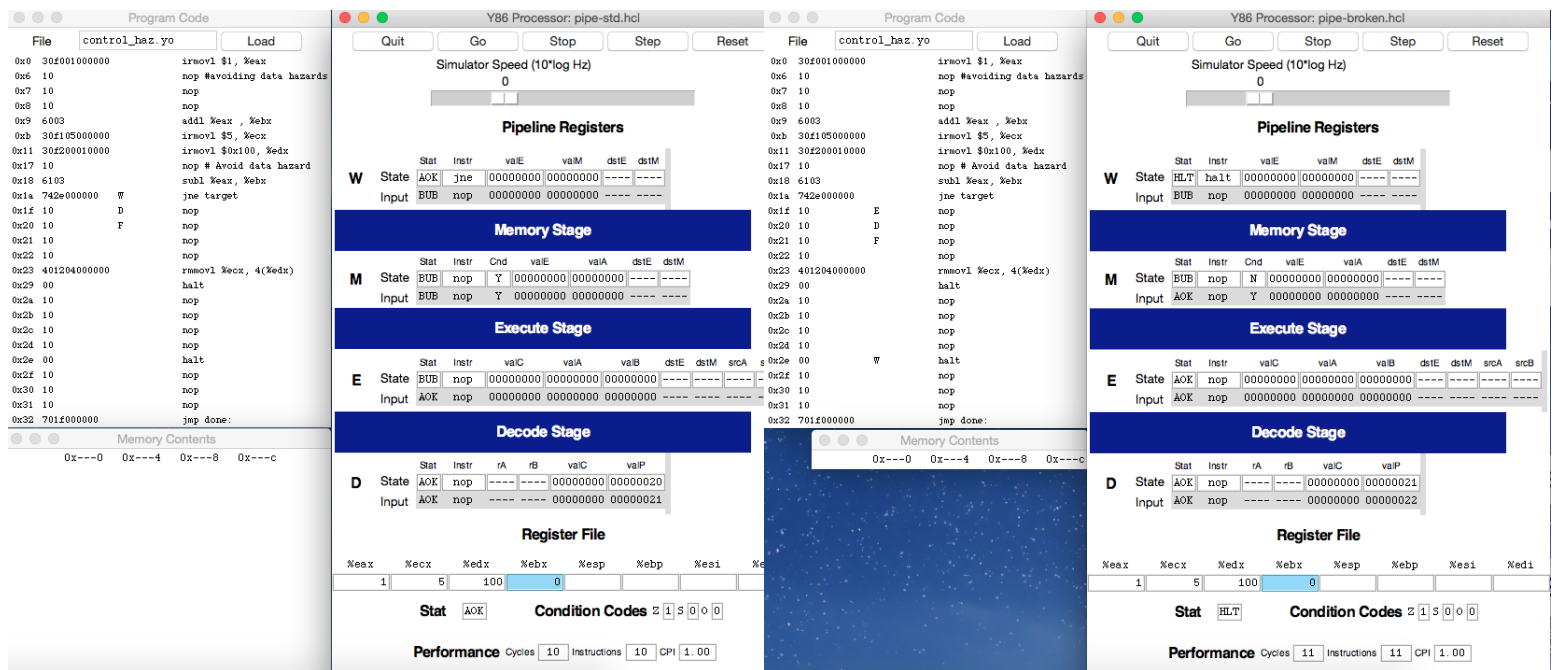
Figure 3: The `jne` instruction reaches the Execute Pipeline Register. It’s now clear that the conditional jump should not have been taken. However, the “target’s” `nop` instruction has just been fetched, and the `halt` instruction has been loaded into the Decode Pipeline Register.



(a) Correct Implementation (std)

(b) Incorrect Implementation (broken)

Figure 4: This step best illustrates the difference between the two versions. The correct implementation cancels the two incorrectly-fetched instructions (by converting them into `nop` bubbles) and then fetches the next instructions from the correct location – the “done” tag. By doing so, the correct version mitigates the control hazard. On the other hand, the incorrect version simply fetches the next instruction from the correct location without cancelling the `halt` instruction or the `nop` instruction, simply progressing them to the Execute and Decode stages respectively.



(a) Correct Implementation (std)

(b) Incorrect Implementation (broken)

Figure 5: You can see the invalid instructions progress through the pipeline stages. You can also see the Correct implementation’s `nop` bubbles move through the stages and the correct instructions getting fetched.

