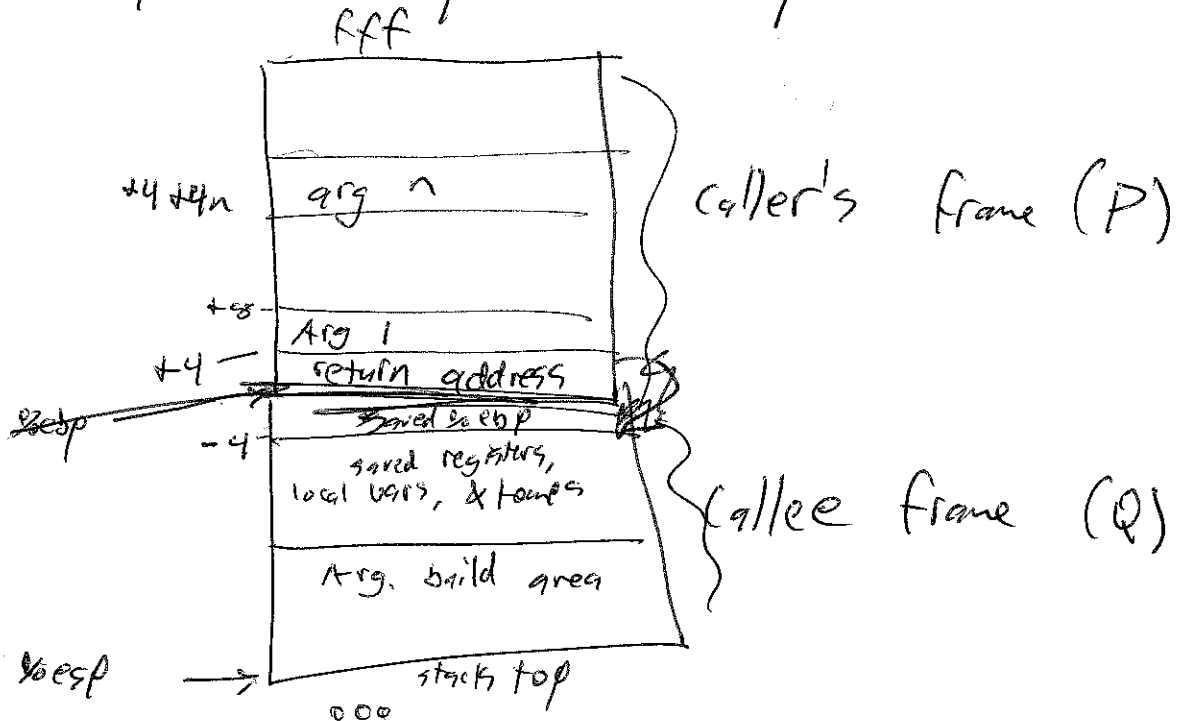


Textbook Notes 3.7-

Procedures

The portion of stack allocated for 1 procedure = stack frame

- %ebp = frame pointer at bottom
- %esp = stack pointer at top



Q also uses the stack for local vars that can't be saved in registers. AAA if:

- not enough registers for all local data
- some local vars are arrays or structures
- the address operator & is applied to a local var (so we need to generate an address for it)

Register use conventions:

- `%eax`, `%edx`, and `%ecx` = caller save
 - ⇒ no need to worry about overwriting them
- `%ebx`, `%esi`, & `%edi` = callee-save
 - ⇒ if you fuck wit them, save 'em first.

So if you're about to call some procedure & you have a val (`y`) you need after, you can

1. store `y` in your stack frame (caller save) OR
2. store `y` in a callee-save register so you know it'll be there after.

- Note GCC always allocates a multiple of 16-bytes for a frame (including 4 bytes for old `%ebp` & 4 bytes for ret address).

Recursion:

- these conventions ⇒ recursion = possible

fact: `#int fact(int n) ⇒ n!`

```
pushl %ebp
movl %esp, %ebp
pushl %ebx          # callee save
subl $4, %esp        # allocate 4 bytes to stack } set up

movl 4(%ebp), %ebx   # set n
movl $1, %ecx        # base return = 1
movl $1, %ebx        # base case? } base of recurse

leal -1(%ebx), %ecx  # compute n-1
movl %ecx, (%esp)    # put n-1 into top of stack } recurse
call fact
movl %eax, %ebx
```

Textbook Notes 3.7 & 3.8 (p.2) & 3.8

```
addl $4 %esp } wrap-up  
popl %ebx  
popl %ebp  
ret
```

Array Allocation & Access:

for data type T & integer constant N ,

$T \quad A[N];$

\Rightarrow 1.) allocates $L \cdot N$ contiguous bytes in mem.
 $L = \text{size (in bytes) of } T$

AND 2.) Makes $A \leftarrow x_0 = \text{pointer to first elem. in mem.}$

- Accessing $A[i] = x_0 + L \cdot i, \quad i \in [0, N-1]$

- There are memory addressing tricks for arrays in IA32

\rightarrow e.g. $\%edx \leftarrow E, \quad \%ecx \leftarrow i$, and $E = \text{array of ints}$

good $\Rightarrow \text{movl } (\%edx, \%ecx, 4), \%eax$

practice $\Rightarrow \text{move } E[i] \text{ into } \%eax$

Pointer arithmetic

$\star(A + i) = A[i]$

\Rightarrow

Nested Arrays:

$\text{int } A[5][3];$

$=$

$\text{typedef int row3_t}[3];$
 $\text{row3_t } A[5];$

Fixed-Size arrays:

define N 16

typedef int fix_matrix[N][N];

⇒ many assembly-level optimizations.

Variable-Size arrays:

- we can do variable-size arrays

```
int val_ele(int n, int A[n][n], int i, int j) {  
    return A[i][j]; }
```

Textbook Notes 3.9 Heterogeneous Data Structures

Structures:

- struct creates a data type that's a group of possibly different data types.
- compiler knows the byte offset of each field
- selection of field occurs solely @ compile time.
machine code has no knowledge of data structures

Unions:

- allow single objects to be referenced according to multiple types.

| | c | i | v | size |
|--|----|---|---|------|
| union U { char c; int i[2]; double v; }; | 53 | 0 | 4 | 12 |
| ⇒ | 43 | 0 | 0 | 0 |
| | | | | 8 |

};

refer to the same memory block. i.e. U can have one of the above fields.

- good for saving space, but can lead to nasty bugs.
e.g. of data type with mutually exclusive fields: (bin. tree)

```
union NODE_U {
```

```
  struct {
```

```
    union NODE_U *left;
```

```
    union NODE_U *right;
```

```
  } internal;
```

```
  double data;
```

```
};
```

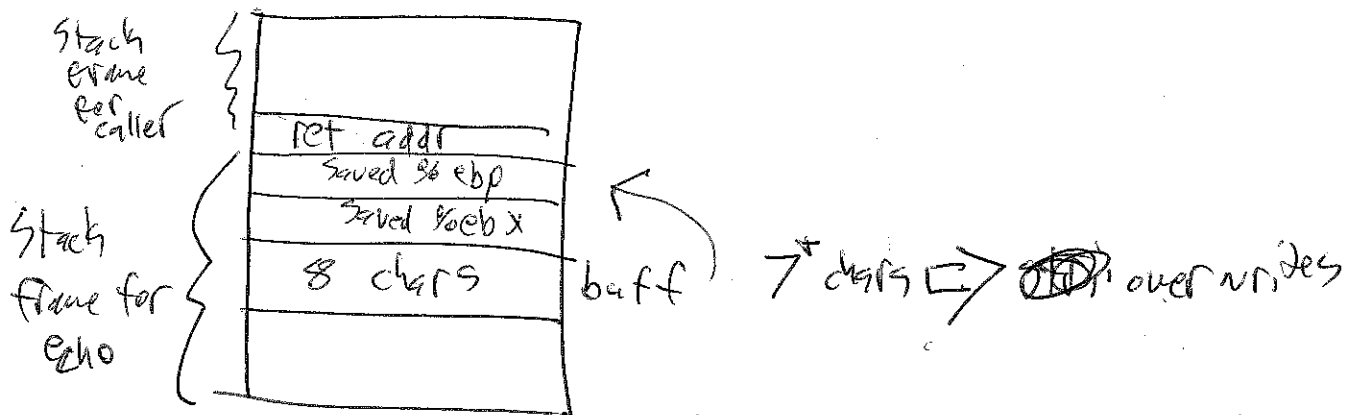
- You can do some cool metaprogramming things w/ unions
- Some systems place restrictions on addressing (requiring each address start at a multiple of 2, 4, or 8).
 - Also data alignment \Rightarrow major improvement.
 - Malloc has to return a pointer that'll work on the worst-case machine (required 4 or 8 ~~alignment~~ alignment)

3.12 Out-of-Bounds mem. references & Buffer Overflow

- Since C doesn't ~~bound-check~~ pointers & local vars are stored on the stack w/ state info, we can get bad program errors

eg. Buffer overflow:

- ~~string~~ array allocated on stack, but size of string exceeds space allocated for the array



\Rightarrow bad %edx | bad %ebp = can't reference local vars or params | bad ret addr = jumps to unexpected location in the program.

Textbook notes

3.12 p. 2

- gets, strcpy, strcat, & sprintf can all generate a byte sequence w/o being given any indication of the size of destination buffer
- Or you can feed the program a string w/ the byte encoding of some executable code (exploit code) + some extra bytes that overwrite the return address w/ a pointer to the exploit code.
 - eg call to start up a shell program
 - or do something, restore code, and return like normal...
- worm = prog. that can run on its own & propagate to other machines
- virus = add-on to other programs (including OSs).

Defense against the dark Arts:

1. Stack Randomization

- changing the memory location of the stack from one program^{run} to another
- Before it was uniform \Rightarrow one attack works on many machines \Rightarrow security monoculture.
- works by allocating random amount of space 0-n bytes on the stack at the start of a program.
- eg by alloc
- \Rightarrow this space not used by the program, but all locations are offset by magic #.
- eg test w/ my OS

- This has become standard practice in linux systems
- part of a larger class of techniques: address-space layout randomization
 - ⇒ entirely different mapping of memory locations each running
- But, ~~an~~ attacker can overcome by brute force attempts w/ different addrs.
- 1 trick is to include a long sequence of "nops" before the exploit code (AKA a "nop sled")
 - ⇒ if the prog lands anywhere in there it just slides to the exploit code

2. Stack Corruption Detection:

- Include a "randomly generated 'canary'" b/t local buffer & the rest of the stack state.
- Before restoring register state, verify canary ~~value~~ is what it should be. If not, abort w/ error.
- GCC does this by default but can be overridden

3. Limiting Executable Code Regions

- ~~do~~ don't give execute permissions to the stack
- But some programs have to dynamically generate & execute code. eg ~~using~~ "just-in-time compilation" techniques for interpreted languages (like Java)

Textbook Notes 6.0

- In reality, memory systems = hierarchy of storage devices w/ different capacities, costs, & and access times.
- CPU Registers hold most frequently used memory
- cache memories = small, fast, nearby CPU, staying close for some memory
- main mem = large, slow disks.
- external mem = disks or tapes of other machines connected by networks.
- Registers = 0 cycles
- cache = 1 - 30 cycles
- main mem = 50 - 200 cycles
- disk/ external = 10's of millions of cycles
- \Rightarrow write programs w/ good locality!
 - frequently use closer memory

6.1.4 Storage Technology Trends

- Price & performance trade-offs
 - SRAM faster than DRAM faster than disk
 - SRAM costs more than DRAM costs more than disk
- But price & performance properties of different storage technologies are changing at very different rates.
- Since 1980, cost & performance of SRAM have improved at about the same rate.
 - access \downarrow by factor of 200
 - cost/byte \downarrow by factor of 300

- But not so for DRAM

- cost/req \downarrow by factor of 130,000!

- access time \downarrow by factor of only 10...

- Disk = even more so!

- cost/req \downarrow by factor of 1,000,000

- access time \downarrow by only factor of 30...

- CPU cycle time \uparrow by 2500x

- effective cycle time divided by # of ~~core~~ processor (cores) \uparrow by 10,000x

- Gap b/t CPU performance & DRAM performance is widening.

Text books Notes 6.2 Locality

- Locality = when programs reference data items that are near other recently referenced data items, or that were recently referenced themselves

⇒ HUGE impact on design & performance

- temporal locality = if it's ref'd once, it's likely to be ref'd again multiple times in near future.
- spatial locality = if it's ref'd, nearby memory will likely be ref'd soon
- Hardware locality = cached memory

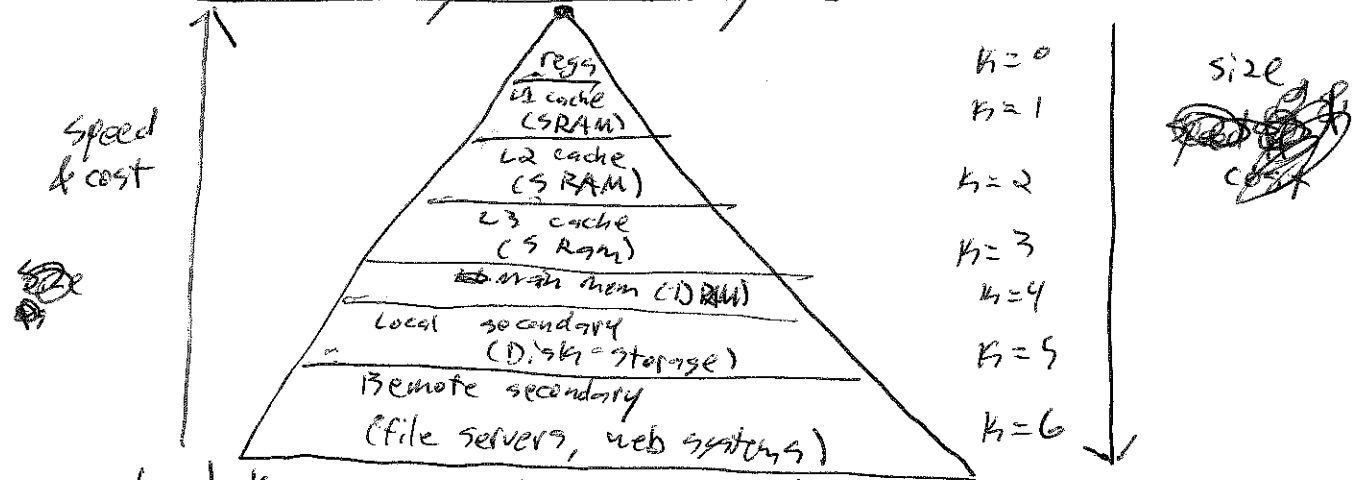
Locality in data ^{access} reference:

- Consider ~~a~~ looping through each elem. in array for ~~searching~~
 - this is "stride-1 reference pattern" or "sequential ref."
 - skipping to every k^{th} elem. ⇒ stride- k reference pattern.
 - $\uparrow k \Rightarrow \downarrow$ spatial locality
- This is big for multi-dimensional arrays
 - access all elems. of 1 row then next row ⇒ good. Stride-1
 - AKA row-major order

Locality of instruction fetches:

- For loop = good spatial locality
 - 'Cause instr.s held in sequential memory blocks
- Also good temporal locality
 - 'Cause instr.s ~~are~~ repeated.
- Smaller loop body & \uparrow iterations ⇒ \uparrow locality

6.3 The memory Hierarchy:



- each level k in hierarchy caches data for level $k+1$
 - contains a subset of (usually) fixed size data blocks from $k+1$
- block size for $L1 \rightarrow L0$ usually ≈ 1 byte
- for $L2 \rightarrow L1$, $L3 \rightarrow L2$, & $L4 \rightarrow L3$ usually $\approx 8-16$ bytes
- for $L5 \rightarrow L4 \sim 100$'s or 1000 's of bytes
- cache hit = when we need d from level $k+1$ & it's stored in level k
- if not, cache miss. $\Rightarrow k$ fetches & stores the data from $k+1$
 - possibly overwriting a block if cache is full.
- empty cache = cold cache. \Rightarrow compulsory / cold misses.
- when there's a miss, cache needs a placement policy. (where to place blocks)
 - most flexible = let the blocks be stored anywhere.
 - this = too expensive for high-level caches
 - \Rightarrow restrictive placement policies
 - but restrictive policies \Rightarrow conflict miss.
 - Also, capacity misses, when each phase has some reasonably constant set of cache blocks (the working set) that's too big to fit in the cache.

Textbook Notes 6.3 and 6.4

- This is why temporal & spatial locality work.
temporal 'cause recent access \rightarrow it's still in cache
spatial 'cause if it's close in mem, the block might include it in cache.

6.4 Cache memories

- L1 cache $\sim 2-4$ clock cycles
- as gap b/t CPU & main mem. grows, more levels of cache
 \Rightarrow L2 ~ 10 clocks, L3 $\sim 30-40$ clock cycles

consider computer system w/ m -bit memory addresses.

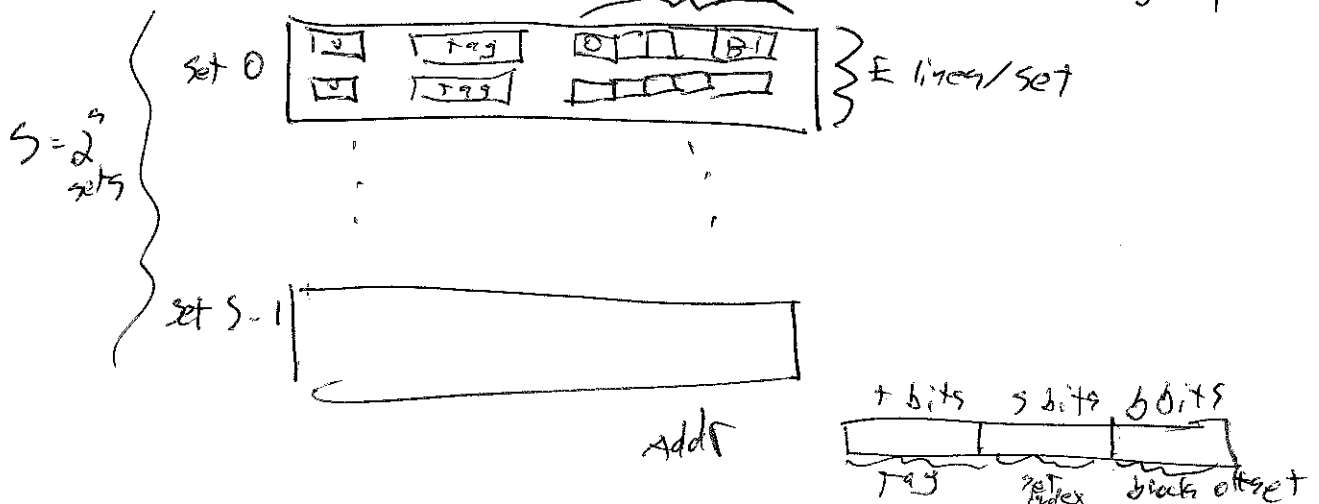
$\Rightarrow M = 2^m$ unique addresses

- A cache organized to $S = 2^s$ cache sets
- Each set = E cache lines
- Each line = block of $B = 2^b$ bytes & valid bit,
& $t = m - (b + s)$ tag bits that uniquely identify the block stored in the cache line.

$\Rightarrow (S, E, B, m)$

$\Rightarrow C = \text{size of cache} = S \times E \times B$
 $B = 2^b \text{ bytes}$

$b = 4$



Direct-Mapped Caches

- $E = 1$ (1 line per set)

- conflict misses \Rightarrow thrashing!

for ($i = 0, j = 0; i < N, j < M; i++, j++$) {

Sum += ~~a~~ $a[i] * b[j]$

$\Rightarrow a[0] \rightarrow a[3]$ cached, then overwritten by $b[0] \rightarrow b[3]$

\Rightarrow overwritten by $a[1] \rightarrow a[4] \Rightarrow$ etc...

- This is especially important for arrays with size 2^k .

- To fix, add an empty buffer after array \Rightarrow the arrays map to different cache sets.

Set Associative:

- $E > 1 \Rightarrow$ tag bit selects the line.

- Line replacement gets weird.

• Fully Associative:

$$S = 1, E = C/B$$

- This is computationally expensive AF

Cache writing is more complicated:

- write through = updating level updates all the lower ones.

\Rightarrow huge bus w/ each write.

- write back = only write it to the next level when that line in the cache is replaced.

\Rightarrow cache must have "dirty bit" that indicates if that line has been modified.

Textbook Notes §.4 ctd

- what about write misses?

"write-allocate" \Rightarrow load block from lower level, then ~~allocate~~ ^{write}
"no-write-allocate" \Rightarrow go directly to lower level.

- write through caches = no write allocate

- write-back caches = write-allocate.

\uparrow
again miss style when programming

- real world:

• I-caches & D-caches = separate (and I = read only).

- Measuring Cache Performance:

1. Miss rate ($\# \text{misses} / \# \text{references}$)

2. Hit rate ($1 - \text{miss rate}$)

3. Hit time (time to deliver a word in cache to CPU)

4. Miss penalty (additional time after a miss)

\Rightarrow

- Trade offs:

\uparrow cache $\Rightarrow \uparrow$ hit rate, \uparrow hit time

\uparrow Block size $\Rightarrow \uparrow$ hit rate for spatial locality, but less lines ^{per set}
 $\Rightarrow \downarrow$ hit rate for temporal locality.

• Modern caches = 32 - 64 bytes

$\Rightarrow \uparrow$ miss penalty

\uparrow Associativity ($\uparrow E$) $\Rightarrow \downarrow$ misses, \uparrow cost, \downarrow speed, \uparrow miss penalty

- write through = simpler & use \uparrow write buffer $\Rightarrow \downarrow$ miss penalty

- write back \Rightarrow less transfers (more important at low levels)

6.5 writing cache-friendly code

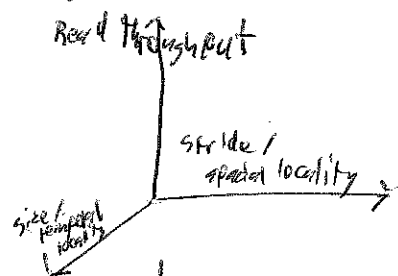
- Good locality \Rightarrow \downarrow miss rate \Rightarrow faster
- Guidelines:

1. make the common case go fast.
2. minimize # of cache misses in each inner loop

- repeated references to local vars are good 'cause compiler can cache them in a reg. file (temporal locality)
- stride-1 ref. patterns are good 'cause codes at all lvs of mem store data as contiguous blocks (spatial locality)

6.6 Putting it Together

- rate a prog. reads ~~mem~~ data from mem = read throughput or read bandwidth. (MB/s)
- Dig memory mountain code (p. 622)
- \uparrow size $\rightarrow \uparrow$ temporal locality
- \uparrow stride $\rightarrow \uparrow$ spatial locality



- each computer's memory mountain is unique!

- eg core i7's

- 4 ridges correspond to regions of temporal locality where entire set fits in L1-3 cache & main mem.
- full order of magnitude b/w ^{highest} L1 ridge & lowest main mem.
- try to exploit locality by getting your program to walk in the peaks, not the valleys

Textbook Notes 6.6 ctd.

~~Class Notes~~ 5/20

Rearranging Loops to ↑ spatial Locality

- consider multiplying 2 $n \times n$ matrices

$$C = AB$$

$$\begin{bmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{bmatrix} = \begin{bmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{bmatrix} \begin{bmatrix} b_{11} & b_{12} \\ b_{21} & b_{22} \end{bmatrix}$$

$$\Rightarrow C_{11} = a_{11}b_{11} + a_{12}b_{21}$$

$$C_{12} = a_{11}b_{12} + a_{12}b_{22}$$

$$C_{21} = a_{21}b_{11} + a_{22}b_{21}$$

$$C_{22} = a_{21}b_{12} + a_{22}b_{22}$$

\Rightarrow 6 functionally equivalent ways to code it

eg i, j, k

```
for (i=0; i<n; i++)  
  for (j=0; j<n; j++) {  
    sum = 0.0;  
    for (k=0; k<n; k++)  
      sum += A[i][k] * B[k][j];  
    C[i][j] += sum;  
  }
```

Results:

- miss rate = better predictor of performance than total # of memory accesses.
- Good locality \Rightarrow constant speed at high levels of n
 - cause intel programmers sweat as shit
 - \Rightarrow mad stride-1 optimization

In general:

1. focus on inner loops (bulk of computation & memory accesses)
2. maximize spatial locality (read data objects sequentially w/ stride k).
3. maximize temporal locality by using data object as often as possible after reading it from mem.

Textbook Notes 4.5 Pipelined Y86

- let's pipeline by adapting SEQ
- first change computation of the PC into the fetch stage

4.5.1 Rearranging the Computation stages (SEQ+)

- PC updates gotta come at the start, "SEQ+"
- make it compute one PC val for the current instruction.
- just circuit rethinking

4.5.2 Inserting Pipeline Registers

- Add registers b/t one stages

1|F holds predicted val of prog. counter

2|D b/t fetch & decode. holds info about most recently fetched instr. for decode

3|E b/t decode & execute. Info about recently decoded instr. for ~~pre~~execute

4|M b/t execute & memory. Info about recently executed instr. for mem stage & branch conditions & branch targets for conditional jumps.

5|W b/t mem & feed back path supplying computed results to reg. file for writing & return addr. to the PC select logic w/ ret. instr.

4.5.3 Rearranging & Relabelling signals

- need multiple copies of signals like valC, srcA, & valE.
- M_stat = status ^{signal} ~~that~~ coming ~~at~~ stored at M pipeline reg.
- m_stat = status signal generated by memory stages control logic blocks.

- note we merge ~~val~~ into val A for pipeline reg. E by choosing either val P from reg. D or val read from A port of reg. file.
- reduces amount of state carried to regs E & M
- we also ~~also~~ associate a status code w/ each instr. for exceptional events.

New PC Instructions:

- we want to get the next PC right after calculating the current one (could \Rightarrow throughput of 1 instr./cycle)
- But if fetched instr. is conditional branch we don't know whether or not to take it till after instr. taken execute stage.
- If instr. == ret, we don't know where to go until instr. passes taken M stage.
- Besides these,
 - for call & jmp, it's val C
 - for others it's val P.
 - we can predict whether conditional jump will be taken ($PC \leftarrow \text{val C}$) or not ($PC \leftarrow \text{val P}$), but we need to deal with the case in which we're wrong
- This = branch prediction.
 - lots of wild systems
 - we'll assume conditionals are always taken... ($PC \leftarrow \text{val C}$)

Textbook Notes Ch. 4.5 p. 2

- $\text{hit} \text{ rate} = \text{"always taken"} = 60\% \text{ success}$
- $\text{"never taken"} = 40\%$
- $\text{"backward taken, forward not-taken"} (\text{BT/NT}) = 65\% \text{ success}$
(think about loops.)
- w/ return, we just don't process any more instrs until ret instr. passes through write-back stage.

4.5.5 Pipeline Hazards

- data hazards & control hazards

eg Prog 1

0x000 irmovl \$10, %edx

0x006 irmovl \$3, %eax

0x00C nop

0x00d nop

0x00e nop

0x00f addl %edx

| 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 |
|---|---|---|---|---|---|---|---|---|----|
| F | D | E | M | W | | | | | |
| | F | D | E | M | W | | | | |
| | | F | D | E | M | W | | | |
| | | | F | D | E | M | W | | |
| | | | | F | D | E | M | W | |
| | | | | | F | D | E | M | W |
| | | | | | | F | D | E | M |
| | | | | | | | F | D | E |

every addl. decodes, %edx & %eax written

Prog 3

0x000 irmovl \$10, %edx

0x006 irmovl \$3, %eax

0x00C nop

0x00d addl %edx, %eax

| | | | | |
|---|---|---|---|---|
| F | D | E | M | W |
| | F | D | E | M |
| | | F | D | E |
| | | | F | D |

when addl. decodes, neither %edx or %eax written yet

⇒ wrong result.

Enumerating classes of data hazards:

1. prog. regs — cause reg. file read in 1 stage, written in another
 2. Prog. counter — conflicts b/t updating & reading prog. counter
 \Rightarrow control hazards. \Rightarrow prediction \Rightarrow mispredicted branches
 / not instruction "exception handling"
 3. Memory — reads/writes occur in 1 stage \Rightarrow no conflict
 (unless self-modifying code...)
 4. CC reg. — written in execut, ~~write~~ read in execute/bench
 \Rightarrow no hazards
 5. Status reg — enables exception handling
- \Rightarrow overall, we only need to worry about reg. data hazards, control hazards, & exception handling.

4.5.6 Stalling

- Holds back instr.s in the decode stage until ~~into~~ ^{generally} its source operands have passed through write-backs.
- Stalling = holding back ~~all~~ subsequent instructions 'till others go through & adding a bubble
 'like a dynamically generated n^o instr.

| | eg | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 |
|-------|-----------------|---|---|---|---|---|---|---|---|---|----|----|
| 0x000 | movl \$10, %edx | F | D | E | M | W | | | | | | |
| 0x006 | imovl \$3, %eax | | F | D | E | M | W | | | | | |
| | bubble | | | | | F | E | M | W | | | |
| | bubble | | | | | | | F | E | M | W | |
| | bubble | | | | | | | | F | E | M | W |
| 0x00C | addl %edx, %eax | | | F | D | D | D | D | E | M | W | |
| 0x00E | hlt | | | | F | F | E | F | D | E | M | W |

- very representation, meh performance.

Text book Notes 4.5. p. 3

4.5.7 Forwarding

- Instead until write is complete, it can pass the value about to be written to pipeline reg. D as the source operand.
- or even the val M_valE to pipeline reg. D
- this passing of result value directly from 1 pipeline stage to an earlier one is "data forwarding"
- avoids stalling, but additional data connections & control logic
- Can be used for values generated by ALU & destined for write port E.
- or for values read from mem. & destined for write port M.
- ALU from: (e_valE, m_valM, M_valE, W_valM, & W_valE)
to: (valA, valB).
- see figs 4.49 to 4.51 for decode-state logic to determine which val to use (from reg. file or a forwarded val?).

Text book Notes 4.5.8 Load/Use Data Hazards

- One class of data hazards can't be handled just by forwarding 'cause mem reads occur late in pipeline.

-eg

when mem read val ~~from~~^{into} %eax write next instr

(eg addl) ~~use~~ needs this as a source operand

- This can be solved by combining stalling & forwarding

- This is called load interlock

• w/ forwarding & load interlock, all data hazards are solved.

4.5.9 Exception Handling

- Exceptioning instruction = instr. causing exception

- eg halt, invalid addr, & invalid instr.

- here, all instr.s prior to this should complete & none after should effect programmer-visible state.

- In bigger system, need make an exception handler

- In pipelined, it's tricky

1. Can get exceptions triggered by multiple instructions at the same time.

eg halt in fetch stage, invalid mem. addr. in mem stage.

⇒ which do we report? usually one in the furthest stage

2. Instr. is fetched, causes exception, but later is canceled due to mispredicted branch.

Textbook Notes 4.5 (p.4)

3. ~~instr.~~ instr. following excepting ~~instr.~~ ~~instr.~~ ~~alters~~ ~~alters~~ alters
prog-visible state before excepting instr. completes.

- we can avoid all these problems by merging exception-handling logic into the pipeline structure
 - ⇒ stg code in each pipeline register.
- that exception status propagates through pipeline w/ the rest of the info for that instruction until the write-back stage.
- then the control logic stops execution
- upon discovery of exception in Mem or write-back, changes to CC reg. and data memory are disabled.
- This strategy solves all the problems

Textbook Notes 4.5.12

- Cycles per instruction (CPI).

• Avg. # of clock cycles taken to do 1 instr.

• Reciprocal of the avg. throughput (in p/cycles)

$$= \frac{C_i + C_o}{C_i} = 1.0 + \frac{C_o}{C_i}$$

C_i = instr.s, C_o = bubbles, $\Rightarrow \frac{C_o}{C_i}$ = bubbles/instr.

$$= 1.0 + lp + mp + rp$$

lp = load penalty = avg. freq. w/ which bubbles injected while stalling for load / use hazards

mp = mispredicted branch penalty = avg. freq. w/ which bubs injected when cancelling instr.s

rp = return penalty = avg. freq. w/ which bubs injected for ret instr.

| Cause | Name | Instr. freq. | Cond. freq. | Bubs | product |
|---------------|------|--------------|-------------|------|---------|
| Load/use | lp | .25 | .20 | 1 | 0.05 |
| mispredict | mp | .20 | .40 | 2 | 0.16 |
| return | rp | .02 | 1.00 | 3 | 0.06 |
| Total penalty | | | | | 0.27 |

$$\Rightarrow CPI = 1.27$$

- Note mispredicted branches = more than half this penalty

unfringed business

- Multi-cycle instr.s \rightarrow 32 cycles for int. division

- Interacting w/ the mem system - self-modifying code.

- no use of virtual addressing.

- no 2-stage 1-cycle mem. access

Textbook Notes 8.1 exceptions

- Partly hardware, partly OS.
- exception = abrupt change in the control flow in response to some change in the processor's state.
- In I_{curr} (encoded in various bits)
change in state = event
- Event egs:
 - virtual memory page fault
 - arithmetic overflow
 - div. by 0.
 - system timer goes off
 - I/O request completes.

assigned by processor designer

number assigned by OS designer
- event \Rightarrow indirect proc. call (exception) \Rightarrow exception table (jump table)
 \Rightarrow OS subroutine (exception handler)
 \Rightarrow
- 1 of 3 things:
 1. ret. to I_{curr}
 2. ret. to I_{next}
 3. handler aborts the program.
- starting addr. of exception table = in the exception table base register.
- when called, ret. addr. pushed on to stack. & processor state. for restoration. (if control \rightarrow kernel, pushed on to kernel's stack).

Classes of Exceptions

1. Interrupt - signal from I/O device - Async - ret. to I_{next}
2. Trap - Intentional Exception - Sync - ret. to I_{next}
3. Fault - Potentially recoverable err. - Sync - might ret. to I_{curr}
4. Abort - non recoverable err. - Sync - no ret.

Note: `syscall` function \Rightarrow do said function in kernel mode (impl if you're doing things requiring permissions).

In Linux: 256 exceptions

- 0 - 31 = Linux i386 defined. 32 - 255 = OS.
- 0 = div. err = fault. - 14 = page fault. - 11 = seg fault
- 128 = sys call (AREA 0x80)

- seg fault = ref. to undefined area in virtual mem. or tryin' to write to a read-only text seg.
- 1 = exit, 2 = fork, 3 = read, 4 = write, 5 = open, 6 = close, 7 = write(2)
- 8 = execve, 20 = get pid, 37 = kill, 29 = pause (wait for sig)
- 27 = alarm

- sys calls made ^{by} app w/ a trap instr. Called int in (location of ^{kernel} kernel)
- see table in `linux/include/sys/syscall.h`

- programs can use `syscall` function, but it's rare.
- usually use a wrapper function
- these = system-level functions

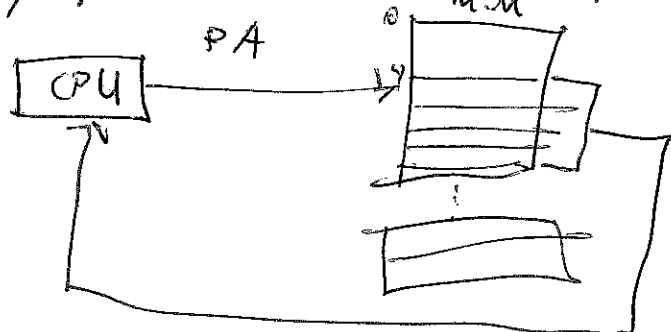
Textbook Notes 9.1

- modern processors use an abstraction of main mem. — virtual mem.
- Avoid mem. conflicts.
- \Rightarrow 3 important benefits:
 1. Uses main memory efficiently (treating it as a cache for an addr. space stored on disks).
 2. Simplifies mem. management (providing each process w/ a uniform addr. space)
 3. Protects addr. space of each proc. from corruption by others.
- virtual mem. = central (all-pervasive)
- virt. mem. (VM) = powerful & dangerous

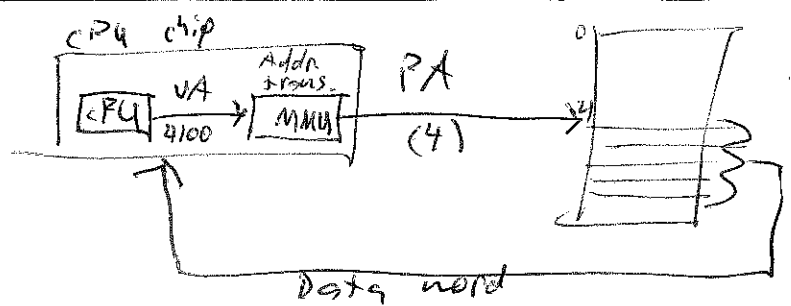
9.1 Physical & virtual addressing

- main mem. = array of M contiguous byte-sized cells.
- each has a physical addr. $(PA)_i \in [0 : (M-1)]$

eg



- w/ ~~VM~~ VM, CPU generates a virtual addr. (VA), which is converted to the right PA before being sent to MM.
- The conversion = addr. translation
- using a memory management unit (MMU)
- using a look-up table stored in main mem.
- contents of which are managed by the OS.



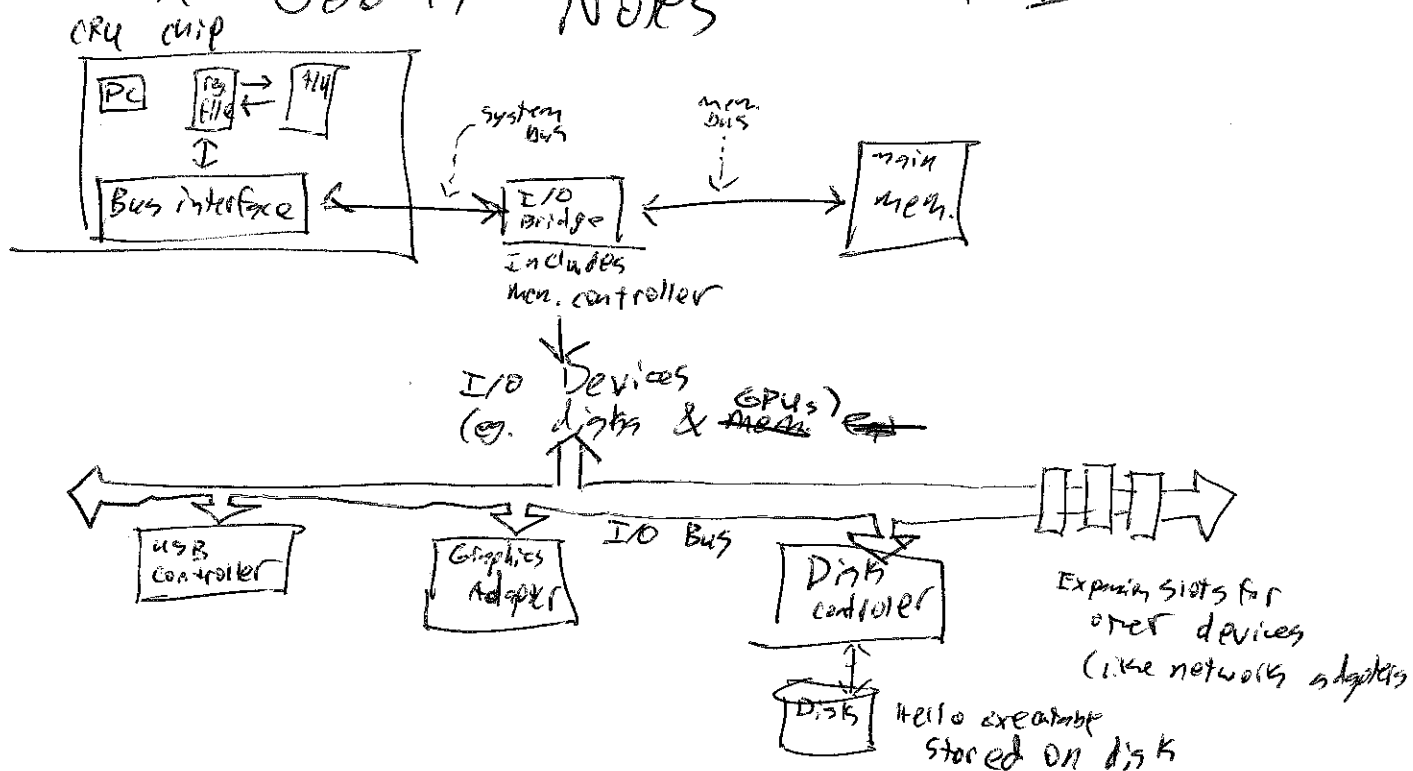
9.2 Addr. Spaces

- an addr. space = ordered seq of nonneg. int. addr.s $\{0, 1, 2, \dots\}$
- if consecutive, it's a linear addr. space.
- In VM, CPU generates virtual addr.s from addr. space of $N = 2^n$ addr.s called "virtual addr. space". $\{0, 1, 2, \dots, N-1\}$
- size of an addr. characterized by # of bits needed to represent the largest addr. \Rightarrow n-bit addr. space (usually 32-bit or 64-bit)
- Also a PA space. M-bytes $\{0, 1, 2, \dots, M-1\}$

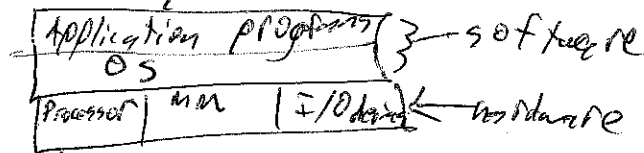
587 Non volatile memory

- DRAM & SRAM = volatile (lose info if supply voltage goes off)
- Non-volatile mem.s retain info
- ROMs
- PROM = 1 write (fuse w/ each mem. cell)
- EPROM = erasable PROM = can be rewritten on order of $10^4 \times$
- EEPROM = electrically EPROM = ^{rewritten} order of 10^5 times
- flash mem. = fast, nonvolatile storage.
 - eg solid state disk (SSD)
- Program stored in ROM = firm ware (eg boot program) & BIOS (Basic I/O system) routines

Text book Notes ch. 1



- CPU = central processing unit = processor
- OS = layer of software b/t application program & the hardware
 1. protects the hardware from misuse by runaway apps.
 2. provide apps w/ simple & uniform mechanisms for manipulating complicated & often wildly different low-level hardware devices



- files = abstractions for I/O devices
- VM = abstraction for both MM & disk I/O devices
- Processes = abstraction for processor, MM, & I/O devices

1.7.3 VM

- code begins at same location for all procs.
- followed by global vars (fixed memory)
- heap = dynamically allocated mem.

